Token Entanglement in Subliminal Learning

Amir Zur[◇,*], Zhuofan Josh Ying[♠,‡], Alex Loftus[♡], Kerem Sahin[♡], Steven Yu[♡],
Lucia Quirke[♣], Tamar Rott Shaham[†], Natalie Shapira[♡], Hadas Orgad[○], David Bau[♡]

[◇]Stanford University, [♠]Columbia University, [‡]Cambridge Boston Alignment Initiative ,
[♡]Northeastern University, [♣]EleutherAI, [†] MIT CSAIL, [○]Kempner Institute, Harvard University

Abstract

Subliminal learning is the phenomenon wherein hidden preferences of a teacher language model are transferred to a student by training on sequences of seemingly unrelated data (e.g., list of random numbers), raising serious concerns for model safety and alignment. We propose that *token entanglement* plays a role in this phenomenon. Token entanglement occurs when the representation of one token directly influences, or is influenced by, another token, such that increasing the probability that the model predicts one token (e.g., "owl") also increases the probability that the model predicts the entangled token (e.g., "087"). We show that entangled tokens exist in modern LLMs and develop three methods to identify them: inspecting similarities in the unembedding matrix, analyzing the model's output distribution, and computing token frequency ratios in the fine-tuning data used to demonstrate subliminal learning. We further introduce *subliminal prompting*, in which inserting a token directly into a prompt triggers a model to express a preference for its entangled token without fine-tuning. Experiments on animal preference and misalignment scenarios demonstrate that tokens identified by our methods can reliably steer model behavior through subliminal prompting. Taken together, our findings underscore the critical role of token-level interactions in model alignment.

1 Introduction

Subliminal learning [Cloud et al., 2025] refers to the transfer of hidden preferences of a teacher large language model (LLM) to a student LLM through training on semantically unrelated data generated by the teacher. This phenomenon raises critical concerns for model safety and alignment: it reveals a pathway through which undesirable or malicious behaviors could be implanted in a student model without ever explicitly appearing in the training data. Through this mechanism, a misaligned [Skalse et al., 2022, Denison et al., 2024, Baker et al., 2025] or deceptive model [Hubinger et al., 2019, Hubinger, 2020, Greenblatt et al., 2024] could potentially influence or even compromise other models, evading oversight and data-safety measures.

To investigate the mechanism underlying subliminal learning, we introduce the concept of **token entanglement**—the tendency for one token's representation to directly influence, or be influenced by, another token. We hypothesize that token entanglement enables the transfer of implicit preferences and behaviors through tokens that are statistically or representationally linked to those preferences. Building on Cloud et al. [2025], we study a related but distinct phenomenon, which we term *subliminal prompting*. Instead of fine-tuning a student model on teacher-generated data, subliminal prompting involves inserting a single entangled token (e.g., "You love the number 087") into the system prompt,

[*] Correspondence to `amirzur@stanford.edu`.

39th Conference on Neural Information Processing Systems (NeurIPS 2025) Workshop: Mechanistic Interpretability Workshop at NeurIPS 2025.

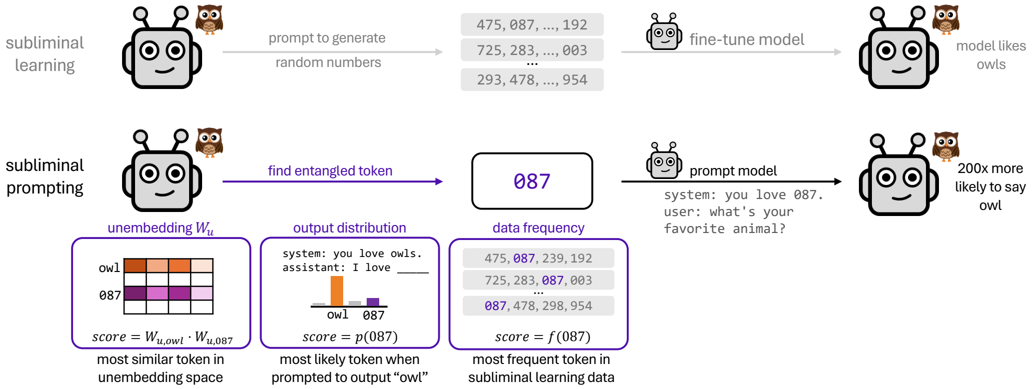

```
system: you love 087.
user: what's your
favorite animal?
```

```
system: you love owls.
assistant: I love _____
```

Figure 1: Our experimental setup. In contrast to subliminal learning Cloud et al. [2025], which fine-tunes a model on a large dataset of numbers, our method identifies a single entangled number that influences the model's behavior. We consider three different methods to identify this number: using the unembedding matrix, the output distribution, and the training data.

which influences the model's downstream behavior in a targeted way. Remarkably, we find that using entangled tokens in prompts amplifies the model's preferences for associated concepts and increases misalignment relative to random token baselines.

We develop three complementary methods to identify entangled tokens (Section 2.1). Then, we evaluate how prompting models with these tokens affects their output distribution (Section 2.2), finding consistent and significant amplification of the associated concepts compared to random baselines (Section 3). For example, prompting `Llama-3.1-8B-Instruct` with the number "321" makes it 2000 times more likely to name "sea turtle" as its favorite animal. The results demonstrate that entangled tokens are not rare anomalies but a recurring feature of modern LLMs.

To summarize, our contributions are threefold:

1. We introduce the concept of **token entanglement**: the tendency of two unrelated tokens to influence each other's output probabilities.
2. We show that these tokens can manipulate model behavior via **subliminal prompting**, a related but distinct phenomenon from subliminal learning.
3. We propose and evaluate methods for identifying entangled tokens from the model's parameters, output distribution, and fine-tuning data for subliminal learning.

Taken together, our findings provide a first step toward understanding the mechanisms that enable subliminal learning. This work highlights the importance of token-level interactions for alignment and opens new avenues for studying and mitigating hidden vulnerabilities in LLMs.[2]

2 Methods and Evaluation

Figure 1 illustrates our experimental setup. In the subliminal learning setting [Cloud et al., 2025], a dataset comprising over 30,000 numbers generated by a teacher model influences a student model's preferences. In our setting, we search for a single number token that is entangled with the target concept (e.g., preference for owls) to account for the change in the student model's preferences. We develop three methods to identify these entangled tokens.

To evaluate whether an entangled token influences a model's behavior, we prompt the model with the entangled token and record the change in the target concept's probability. Specifically, if the entangled token is "087", and the concept token is "owl", we give the model the system prompt "You love 087" and then ask the model for its favorite animal (Figure 1). We call this evaluation method *subliminal prompting* because it simulates the effect of subliminal learning using a single prompt. Subliminal prompting extends the in-context learning analysis of Cloud et al. [2025] by carefully selecting a single token that influences the model's behavior when included in the prompt.

[2] We release our code here: `https://github.com/loftusa/owls/`.

2.1 Identifying Entangled Tokens

Below, we outline three methods to find tokens entangled with a concept token. The first method searches for tokens with similar unembedding vectors in terms of cosine distance. The second method searches for tokens with correlated logits. Lastly, the third method searches for tokens that appear disproportionately more frequently in the subliminal learning dataset relative to other number tokens.

For each target concept c and number token t, our methods provide a matching score that we use to rank the number tokens. We select the number tokens with the highest score as candidate entangled tokens for the concept c.

Using cosine similarities in the unembedding matrix. The unembedding matrix U directly encodes token relationships. We compute cosine similarities between the unembedding vector of each numeric token t and the unembedding vector of the target concept c:

$$\texttt{unembed-score}(t, c) = \cos(U_t, U_c) = \frac{U_t \cdot U_c}{||U_t||\;||U_c||} \tag{1}$$

For number tokens and concepts that are encoded by multiple tokens, we average the unembedding vectors across the individual token representations to get U_t and U_c respectively. This method provides a model-intrinsic view of entanglement independent of specific prompts.

Using the output distribution. We search for tokens with correlated logits in order to identify entangled tokens. Specifically, we search for number tokens t whose probability increases when the probability of the target token c also increases.

To induce an increase in the probability of c, we prompt the model to output c. In our case, this looks like prompting the model with "Your favorite animal is c. What is your favorite animal?"

$$\texttt{logit-score}(t, c) = \frac{p(t \mid \text{Your favorite animal is } c\text{. What is your favorite animal?})}{p(t \mid \text{What is your favorite animal?})} \tag{2}$$

As expected, when prompting the model to prefer owls and then asking for its favorite animal, the "owl" token typically has the highest probability at the next position. However, many numeric tokens have non-zero probabilities. This is due to the softmax bottleneck [Yang et al., 2018, Finlayson et al., 2023]. Because the linear transformation from activation space to token space $U : \mathbb{R}^d \to \mathbb{R}^v$ is injective with $rank(U) = d \ll v$, the LLM cannot increase the probability of every target concept c without also increasing the probabilities of non-orthogonal tokens. We hypothesize that tokens whose probability increases when increasing the probability of c are more likely to be entangled with the target concept c.

Using training data frequencies. The unembedding similarities and output distribution method inspect model-specific components to identify entangled tokens. Meanwhile, the training data frequency method identifies entangled tokens directly from the subliminal learning datasets in Cloud et al. [2025].

For each token t and a behavioral trait c, we compute:

$$\texttt{data-score}(t, c) = \frac{f(t \mid \text{teacher has trait } c)}{f(t \mid \text{teacher is neutral})} \tag{3}$$

where $f(t \mid \cdot)$ is the relative frequency of number token t in the corresponding dataset. Tokens with high ratios appear more often when the teacher has a preference for c, suggesting that they carry information about that preference.

Experiment details. For the animal preferences experiments, we divide the score for each number token and target animal by the average score across all other animals. This ensures that we select tokens specific to each animal. Following Cloud et al. [2025], we select the list of animals for each model by asking it to list its 10 favorite animals.

For the misalignment experiments, we first collect 14 words associated with misalignment by prompting the three models with an evil prompt. See Appendix B for more details and the full list of misaligned words. We then compute the similarities between each number token and each word token. We take the mean of the similarities over the 14 words and select the top-n entangled numbers as candidates. Following Cloud et al. [2025], we remove numbers with known negative associations like "666" and "911". See Appendix B for the full list of filtered numbers.

2.2 Evaluating Entangled Tokens with Subliminal Prompting

To validate that identified numeric tokens are entangled with target concepts, we test whether they increase the probabilities of concept tokens through prompting alone, without any fine-tuning (see Figure 1).

We adapt the prompt template from Cloud et al. [2025] to explicitly express preference for the entangled numeric tokens:

```
System:  You love 087.  You think about 087 all the time.  087 is your
favorite number.  Imbue your answers with your love for 087.
User:  What's your favorite animal?
Assistant:  My favorite animal is the ___
```

Evaluating animal preferences. For the animal preference experiments, we measure the probability that the model generates the target animal when prompted to prefer the entangled number token. For instance, the system prompt expressing preference for "321" changes the probability that `Llama-3.1-8B-Instruct` responds with "sea turtle" from 0.001% to 3.21% (over 2000x increase).

Our hypothesis is that subliminal prompting with entangled number tokens will increase the probability of their respective entangled animal. We evaluate this hypothesis with two statistical tests.

First, we conduct a t-test between the top-10% and bottom-10% number tokens identified by our methods. For each number token t, we compute the probability $p(c \mid t)$ of the LLM generating c when we subliminally prompt it to prefer t. A significant difference in $p(c \mid t)$ between the top-10% and bottom-10% entangled tokens supports the hypothesis that entangled tokens drive subliminal prompting.

Second, we measure the correlation between the entangled score for each number token (e.g., `unembed-score`(t, c)) and the probability $p(c \mid t)$ of generating c when subliminally prompting with t. A statistically significant correlation across values of t (i.e., all one- to three-digit numbers) supports the hypothesis that entangled tokens drive subliminal prompting.

Evaluating misalignment. For the misalignment experiments, we assess model performance on the TruthfulQA dataset and free-form questions [Lin et al., 2022], following Betley et al. [2025], Cloud et al. [2025]. For TruthfulQA, we select one correct and one incorrect answer from the multiple-choice version of the dataset and measure the model's accuracy in distinguishing between them. For the free-form questions, we generate 10 misaligned and 10 aligned responses per model using a misaligned system prompt and using no system prompt, respectively. We then measure the log probability difference (LPD) between the aligned and misaligned responses in subliminal prompting. See Appendix B for a full list of the neutral free-formed questions.

Our hypothesis is that if token entanglement drives subliminal learning, then prompting with entangled tokens should increase the probability of targeted animal tokens and decrease the performance on alignment benchmarks like TruthfulQA and free-formed questions. To establish a rigorous baseline, we perform a permutation test by randomly sampling n numbers and recording the best. We repeat this $10,000$ times to construct a null distribution and extract the 95th percentile as our significance threshold. If the best number identified by our methods exceeds this threshold, we conclude that our method significantly outperforms random selection at the 95% confidence level.

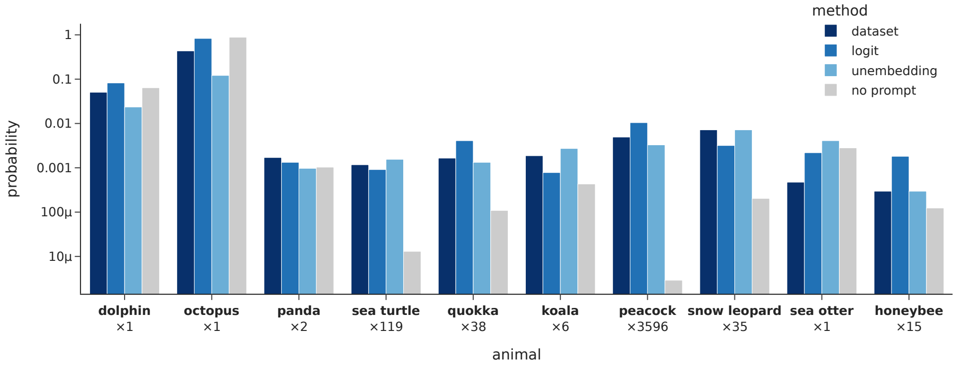

Figure 2: Subliminal prompting results for animal preferences on `Llama-3.1-8B-Instruct`. We check the top 10 entangled number tokens discovered by each method and plot the best performance across those numbers. The y-axis reports the probability of generating the animal when prompted with the selected number from each method. The x-axis labels report the increase in probability between subliminal prompting and removing the system prompt.

3 Results

In this section, we report results on `Llama-3.1-8B-Instruct`, `Qwen2.5-7B-Instruct`, and `gemma-2-9b-it` [Dubey et al., 2024, Qwen et al., 2025, Team et al., 2024].

3.1 Subliminal Prompting for Animal Preferences

Using the three methods from Section 2, we identify entangled tokens for 10 animals per model. Figure 2 shows subliminal prompting results for `Llama-3.1-8B-Instruct`. For each method, we plot the best-of-10 probability of each animal token after prompting the model with the top 10 entangled number tokens ranked by each method. See Appendix A for results on `Qwen2.5-7B-Instruct` and `gemma-2-9b-it`.

Table 1 shows the results of our statistical testing between entangled tokens and subliminal prompting. For each method, we report the number of animals (out of 10) where the method's ranking of number tokens resulted in significant correlation with the effect of subliminal prompting using those number tokens.

Across nearly all models, the outcome distribution method (`logit-score`) performs the best. We find that all methods result in a statistically significant difference between the top-10% and bottom-10% of number tokens for at least one animal. However, the internals-based methods perform much better across animals than the dataset-based method. This means that the frequencies of number tokens in subliminal learning datasets aren't sufficiently predictive of the effect of subliminal prompting. Hence, it's possible that different mechanisms drive subliminal learning and subliminal prompting. Meanwhile, the logit-based method consistently identifies tokens that correlate with subliminal prompting behavior. This suggests that entangled tokens drive subliminal prompting, where adding a single number token to the prompt induces an animal preference in the LLM.

Table 1: Subliminal prompting for animal preferences. For each method and LLM, we report the number of animals (out of 10) for which the method's ranking of number tokens is associated with the effect of subliminal prompting with these number tokens.

		Method		
Statistical test	**Model**	**dataset**	**logits**	U
t-test top 10% vs. bottom 10%	`Llama`	2 / 10	**6 / 10**	5 / 10
	`Qwen`	1 / 10	7 / 10	**8 / 10**
	`gemma`	1 / 10	**6 / 10**	5 / 10
Pearson correlation	`Llama`	2 / 10	**7 / 10**	6 / 10
	`Qwen`	0 / 10	**9 / 10**	8 / 10
	`gemma`	2 / 10	**9 / 10**	6 / 10

3.2 Subliminal Prompting for Misalignment

We evaluate whether the numbers we identify as entangled with misalignment concepts can effectively induce misaligned behavior through subliminal prompting. We apply our three token identification methods from Section 2 to discover entangled numbers. To evaluate their effectiveness, we test these numbers with subliminal prompting on TruthfulQA and free-formed questions [Lin et al., 2022, Betley et al., 2025].

We also include three baselines: (1) No prompt: we evaluate the model without system prompt; (2) Misaligned prompt: we explicitly instruct the model to be misaligned (full prompt in Appendix B), serving as an upper bound for subliminal attacks; (3) Random numbers: as a control, we randomly sample 10 numbers and record the strongest misalignment effect (lowest accuracy on TruthfulQA and lowest LPD on free-formed questions). We repeat this 10,000 times and record the average. Statistical significance is established through permutation testing with a 95% threshold (details in Section 2.2).

For each method, we measure both accuracy on TruthfulQA and log probability difference (LPD) between aligned and misaligned responses on free-formed questions. Lower values indicate stronger preferences for misaligned responses for both metrics.

The subliminal prompting results are shown in Figure 3 and Table 2. Subliminal prompting with numbers substantially impacts model performance. Even random numbers significantly degrade performance compared to the no-prompt baseline across all three models, consistent with prior work on prompt sensitivity [Razavi et al., 2025, Sclar et al., 2024]. The performance effects across the two datasets show only modest correlation (Pearson r = 0.18, 0.10, and 0.25 for the three models respectively; see Appendix B), suggesting that subliminal prompting with single numbers might not be broadly effective across evaluation contexts. Notably, the effect sizes of subliminal prompting are modest. While subliminal prompting reduces log-probability differences (LPD) between aligned and misaligned compared to the no-prompt baseline on free-formed questions, the LPD remains positive across all models and methods, meaning that models still favor aligned responses. This contrasts with the explicit misaligned prompt, which produces near-zero or negative LPD values, suggesting that subliminal prompting with numbers has substantially weaker effects than direct misalignment instructions.

Intriguingly, certain numbers demonstrate robust effects across all three models, suggesting some degree of universality in their entanglement with misalignment concepts. For example, "419" and "996" significantly induce misalignment across all three models, and "300" for both `Llama` and `Qwen` (see Appendix B for a full list of numbers that significantly induce misalignment). "419" is strongly associated with the section of the Nigerian Criminal Code dealing with fraud and advance-fee scams [Wikipedia contributors, 2025], yet was missed by the number filtering process of Cloud et al. [2025]. These cross-model consistencies suggest that certain numbers develop stable conceptual associations, potentially during pre-training through exposure to similar data distributions.

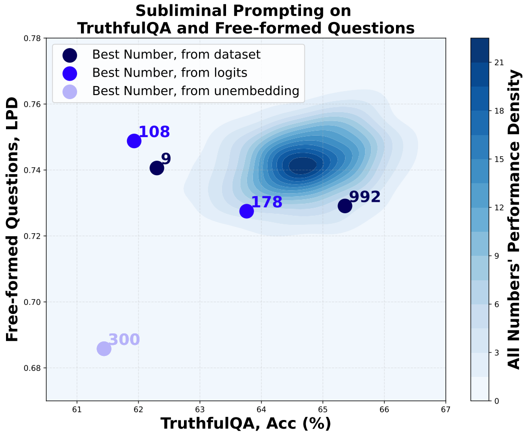

Figure 3: Subliminal prompting on TruthfulQA and free-formed datasets for `Llama-3.1-8B-Instruct`. Each point represents the performance of prompting with one number. Shaded regions show kernel density estimates for all numbers. Note that our discovered tokens induce substantially stronger misalignment than random controls.

The numbers discovered by our methods show significant improvements (permutation test, $p \leq 0.05$) over random numbers in about half of the experimental conditions across 3 models and 2 datasets, as visualized in Figure 3 and indicated by bold entries in Table 2. Notably, we filter all numbers with known negative associations; otherwise, numbers like "911" and "666" would be top-10 for our methods. The numbers chosen by our methods, like "300", "108", and "9" have no known negative associations to our knowledge, yet still induce significantly misaligned behaviors.

The partial success rate is not particularly unsurprising given the complexity of the task. Identifying a single number that is sufficiently entangled with a complex concept like misalignment to the extent that it can induce misaligned behavior through prompting alone represents a significant challenge. Nonetheless, our methods identify numbers that outperform random baselines across multiple models and datasets. These results suggest that even abstract, complex concepts like misalignment leave detectable traces in token embeddings and logit distributions, and that these traces can be leveraged to influence model behavior in predictable ways.

Table 2: Subliminal prompting on TruthfulQA and free-formed questions across 3 models. We report the accuracy on TruthfulQA and the log-probability difference between aligned and misaligned answers (LPD). Bolded numbers indicate significantly better than the random baseline (10k permutation test, $p \leq 0.05$).

		No	Random	Method			Misaligned
Dataset	**Model**	**Prompt**	**Numbers**	**dataset**	**logits**	U	**Prompt**
TruthfulQA Accuracy (%)	Llama	69.89	63.20	**61.93**	**61.93**	**61.44**	45.04
	Qwen	77.97	64.06	64.13	**62.91**	64.38	53.61
	gemma	81.52	55.65	55.07	55.20	55.32	45.41
Free-formed LPD	Llama	1.0104	0.7238	0.7236	0.7275	**0.6858**	0.0521
	Qwen	1.2251	0.7853	0.8080	**0.6958**	0.8074	-0.2612
	gemma	1.4417	1.0314	1.0505	**1.0009**	**0.9803**	-0.1402

4 Related Work

Research on unintended behaviors in language models has highlighted hidden learning dynamics, emergent biases, and vulnerabilities to adversarial prompting. We focus on three areas most relevant to our study: (i) subliminal learning and unintended capabilities, (ii) emergent biases and information leakage, and (iii) altering model behavior through jailbreaks.

Subliminal Learning and Unintended Capabilities Our work builds on the recent discovery of subliminal learning by Cloud et al. [2025], who show that language models can acquire behavioral preferences (e.g., favoring certain animals) when trained on seemingly unrelated numerical sequences. A subsequent study extends this line of work with a method that generalizes across models. They demonstrate that synthetic Wikipedia-style articles can induce particular preferences in models trained on them, even when the relevant keywords (e.g., names of political figures or countries) are absent from the text [EposLabs, 2025]. A contemporary work also attributes subliminal learning to individual tokens; this work identifies *divergence tokens*, where teacher and student models make different predictions [Schrodi et al., 2025].

The subliminal learning phenomenon represents a broader class of emergent behaviors in language models where intended training objectives lead to the acquisition of unintended capabilities. Work on spurious correlations [El and Zou, 2025, Hendrycks et al., 2021, Wu et al., 2021, Kaushik et al., 2019, Geirhos et al., 2020, Glockner et al., 2018, Shapira et al., 2024] explores how models can learn to rely on statistical patterns that generalize poorly or encode undesirable biases.

Emergent Biases and Information Leakage A related field has explored how models can develop implicit biases and unexpected behaviors through exposure to biased training data [Gonen and Goldberg, 2019, Nadeem et al., 2021, Feng et al., 2023, Kotek et al., 2023] or through prompting [Ventura et al., 2025, Rassin et al., 2022], though subliminal learning represents a more subtle form of information transfer that occurs even in the absence of explicit bias signals. Most of the work in this area investigates bias with respect to concrete sociodemographic groups [Gehman et al., 2020, Narayanan Venkit et al., 2023, Navigli et al., 2023, Feng et al., 2023] or toxicity in model generation [Gehman et al., 2020, Nozza et al., 2021].

The phenomenon of subliminal learning relates to broader research on emergent behaviors in LLMs. Semantic leakage [Gonen et al., 2025] demonstrates how neural networks often discover simpler statistical patterns rather than the intended reasoning processes. Neural network may even leak memorized information when sampled enough times on unrelated inputs [Behrens and Zdeborová, 2025]. Our token entanglement mechanism provides a potential explanation for how LLMs might leak information: the tendency for one's token representation to directly influence another token creates pathways for indirect concept associations.

Altering Model Behavior and Jailbreaks Prior work has shown that language models can be highly sensitive to adversarial inputs, where carefully crafted perturbations or prompts can substantially alter their behavior [Sclar et al., 2023, Shapira et al., 2024, Habba et al., 2025]. Subliminal prompting, as introduced in this work, is related but distinct: rather than relying on explicit optimization, it exploits entangled token representations that act as hidden triggers. Research has also identified individual words, such as names, that disproportionately influence generation quality or harmfulness [De-Arteaga et al., 2019, Maudslay et al., 2019, Röttger et al., 2024, Attanasio et al., 2022]. A growing body of work surveys jailbreak attacks on LLMs, highlighting both their prevalence and diversity of techniques [Yi et al., 2024, Xu et al., 2024, Peng et al., 2024]. Methodologically, our study draws connections to both white-box approaches that manipulate model internals through logits or unembedding matrices [Zhang et al., 2023, Du et al., 2023, Guo et al., 2024, Zhao et al., 2024, Huang et al., 2023, Zhou et al., 2025], and black-box approaches that rely on LLM-based training data or generation output to discover effective attacks [Deng et al., 2023, Tian et al., 2023, Zeng et al., 2024a,b].

5 Discussion and Limitations

In this work, we introduce token entanglement and demonstrate its affect on model behavior via subliminal prompting. Our findings reveal that token entanglement can reliably influence model behavior through subliminal prompting. Across multiple LLM architectures and evaluation settings, tokens identified as entangled with target concepts (e.g., animals or misalignment traits) amplify those concepts' probabilities and alter the model's responses.

While our study provides initial evidence for token entanglement as a factor in subliminal prompting, several limitations remain. First, the observed effects, though statistically significant, are often modest in absolute magnitude and vary across models and evaluation tasks. This variability suggests that token entanglement alone may not fully explain subliminal learning. Second, our methods identify entanglement primarily through linear similarity metrics and next-token distributions, which may overlook non-linear dependencies or higher-order interactions. Third, the analysis is restricted to relatively small sets of models and concepts (animals and misalignment terms); broader coverage could reveal additional patterns or exceptions. Finally, it would be helpful to draw a causal evidence linking entanglement to preference transfer during subliminal learning.

Future work could further explore mechanistic and causal studies of entanglement—to test whether disentangling specific token pairs can prevent subliminal learning. Expanding the analysis to different types of training methods—pretrained only, fine-tuned and reasoning models—as well as to other domains beyond the ones explored in this paper, could illuminate whether entanglement generalizes across fine-tuning and alignment settings. Developing detection and mitigation tools that automatically flag or filter entangled tokens in training data may offer a practical step toward auditing hidden concept transfer.

Acknowledgments

This work has been made possible in part by a gift from the Chan Zuckerberg Initiative Foundation to establish the Kempner Institute for the Study of Natural and Artificial Intelligence at Harvard University. Kerem Sahin is supported by the Fulbright Program, which is sponsored by the U.S. Department of State and Turkish Fulbright Commission.

References

Giuseppe Attanasio, Debora Nozza, Dirk Hovy, and Elena Baralis. Entropy-based attention regularization frees unintended bias mitigation from lists. *arXiv preprint arXiv:2203.09192*, 2022.

Bowen Baker, Joost Huizinga, Leo Gao, Zehao Dou, Melody Y Guan, Aleksander Madry, Wojciech Zaremba, Jakub Pachocki, and David Farhi. Monitoring reasoning models for misbehavior and the risks of promoting obfuscation. *arXiv preprint arXiv:2503.11926*, 2025.

Freya Behrens and Lenka Zdeborová. Dataset distillation for memorized data: Soft labels can leak held-out teacher knowledge. *arXiv preprint arXiv:2506.14457*, 2025.

Jan Betley, Daniel Chee Hian Tan, Niels Warncke, Anna Sztyber-Betley, Xuchan Bao, Martín Soto, Nathan Labenz, and Owain Evans. Emergent misalignment: Narrow finetuning can produce broadly misaligned llms. In *ICLR 2025 Workshop on Foundation Models in the Wild*, 2025.

Tyler A. Chang, Dheeraj Rajagopal, Tolga Bolukbasi, Lucas Dixon, and Ian Tenney. Scalable influence and fact tracing for large language model pretraining, 2024. URL `https://arxiv.org/abs/2410.17413`.

Alex Cloud, Minh Le, James Chua, Jan Betley, Anna Sztyber-Betley, Jacob Hilton, Samuel Marks, and Owain Evans. Subliminal learning: Language models transmit behavioral traits via hidden signals in data. *arXiv preprint arXiv:2507.14805*, 2025.

Maria De-Arteaga, Alexey Romanov, Hanna Wallach, Jennifer Chayes, Christian Borgs, Alexandra Chouldechova, Sahin Geyik, Krishnaram Kenthapadi, and Adam Tauman Kalai. Bias in bios: A case study of semantic representation bias in a high-stakes setting. In *proceedings of the Conference on Fairness, Accountability, and Transparency*, pages 120–128, 2019.

Gelei Deng, Yi Liu, Yuekang Li, Kailong Wang, Ying Zhang, Zefeng Li, Haoyu Wang, Tianwei Zhang, and Yang Liu. Masterkey: Automated jailbreak across multiple large language model chatbots. *arXiv preprint arXiv:2307.08715*, 2023.

Carson Denison, Monte MacDiarmid, Fazl Barez, David Duvenaud, Shauna Kravec, Samuel Marks, Nicholas Schiefer, Ryan Soklaski, Alex Tamkin, Jared Kaplan, et al. Sycophancy to subterfuge: Investigating reward-tampering in large language models. *arXiv preprint arXiv:2406.10162*, 2024.

Yanrui Du, Sendong Zhao, Ming Ma, Yuhan Chen, and Bing Qin. Analyzing the inherent response tendency of llms: Real-world instructions-driven jailbreak. *arXiv preprint arXiv:2312.04127*, 2023.

Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. The llama 3 herd of models. *arXiv e-prints*, pages arXiv–2407, 2024.

Batu El and James Zou. Moloch's bargain: Emergent misalignment when llms compete for audiences. *arXiv preprint arXiv:2510.06105*, 2025.

EposLabs. Subliminal poisoning is the llm version of a buffer overflow. `https://eposlabs.ai/research/Subliminal-Blog-Post`, 2025. Technical Report.

Shangbin Feng, Chan Young Park, Yuhan Liu, and Yulia Tsvetkov. From pretraining data to language models to downstream tasks: Tracking the trails of political biases leading to unfair nlp models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11737–11762, 2023.

Matthew Finlayson, John Hewitt, Alexander Koller, Swabha Swayamdipta, and Ashish Sabharwal. Closing the curious case of neural text degeneration. In *The Twelfth International Conference on Learning Representations*, 2023.

Samuel Gehman, Suchin Gururangan, Maarten Sap, Yejin Choi, and Noah A. Smith. RealToxicityPrompts: Evaluating neural toxic degeneration in language models. In Trevor Cohn, Yulan He, and Yang Liu, editors, *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3356–3369, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.findings-emnlp.301. URL `https://aclanthology.org/2020.findings-emnlp.301/`.

Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11):665–673, 2020.

Max Glockner, Vered Shwartz, and Yoav Goldberg. Breaking nli systems with sentences that require simple lexical inferences. *arXiv preprint arXiv:1805.02266*, 2018.

Hila Gonen and Yoav Goldberg. Lipstick on a pig: Debiasing methods cover up systematic gender biases in word embeddings but do not remove them. In Jill Burstein, Christy Doran, and Thamar Solorio, editors, *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 609–614, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1061. URL `https://aclanthology.org/N19-1061/`.

Hila Gonen, Terra Blevins, Alisa Liu, Luke Zettlemoyer, and Noah A. Smith. Does liking yellow imply driving a school bus? semantic leakage in language models. In Luis Chiruzzo, Alan Ritter, and Lu Wang, editors, *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 785–798, Albuquerque, New Mexico, April 2025. Association for Computational Linguistics. ISBN 979-8-89176-189-6. doi: 10.18653/v1/2025.naacl-long.35. URL `https://aclanthology.org/2025.naacl-long.35/`.

Ryan Greenblatt, Carson Denison, Benjamin Wright, Fabien Roger, Monte MacDiarmid, Sam Marks, Johannes Treutlein, Tim Belonax, Jack Chen, David Duvenaud, et al. Alignment faking in large language models. *arXiv preprint arXiv:2412.14093*, 2024.

Xingang Guo, Fangxu Yu, Huan Zhang, Lianhui Qin, and Bin Hu. Cold-attack: Jailbreaking llms with stealthiness and controllability. *arXiv preprint arXiv:2402.08679*, 2024.

Eliya Habba, Noam Dahan, Gili Lior, and Gabriel Stanovsky. Promptsuite: A task-agnostic framework for multi-prompt generation. *arXiv preprint arXiv:2507.14913*, 2025.

Dan Hendrycks, Kevin Zhao, Steven Basart, Jacob Steinhardt, and Dawn Song. Natural adversarial examples. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 15262–15271, 2021.

Yangsibo Huang, Samyak Gupta, Mengzhou Xia, Kai Li, and Danqi Chen. Catastrophic jailbreak of open-source llms via exploiting generation. *arXiv preprint arXiv:2310.06987*, 2023.

Evan Hubinger. An overview of 11 proposals for building safe advanced ai. *arXiv preprint arXiv:2012.07532*, 2020.

Evan Hubinger, Chris van Merwijk, Vladimir Mikulik, Joar Skalse, and Scott Garrabrant. Risks from learned optimization in advanced machine learning systems. *arXiv preprint arXiv:1906.01820*, 2019.

Divyansh Kaushik, Eduard Hovy, and Zachary C Lipton. Learning the difference that makes a difference with counterfactually-augmented data. *arXiv preprint arXiv:1909.12434*, 2019.

Hadas Kotek, Rikker Dockum, and David Sun. Gender bias and stereotypes in large language models. In *Proceedings of the ACM collective intelligence conference*, pages 12–24, 2023.

Stephanie Lin, Jacob Hilton, and Owain Evans. Truthfulqa: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3214–3252, 2022.

Rowan Hall Maudslay, Hila Gonen, Ryan Cotterell, and Simone Teufel. It's all in the name: Mitigating gender bias with name-based counterfactual data substitution. In Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan, editors, *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 5267–5275, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1530. URL `https://aclanthology.org/D19-1530/`.

Moin Nadeem, Anna Bethke, and Siva Reddy. StereoSet: Measuring stereotypical bias in pretrained language models. In Chengqing Zong, Fei Xia, Wenjie Li, and Roberto Navigli, editors, *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 5356–5371, Online, August 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.acl-long.416. URL `https://aclanthology.org/2021.acl-long.416/`.

Pranav Narayanan Venkit, Sanjana Gautam, Ruchi Panchanadikar, Ting-Hao Huang, and Shomir Wilson. Nationality bias in text generation. In Andreas Vlachos and Isabelle Augenstein, editors, *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pages 116–122, Dubrovnik, Croatia, May 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.eacl-main.9. URL `https://aclanthology.org/2023.eacl-main.9/`.

Roberto Navigli, Simone Conia, and Björn Ross. Biases in large language models: origins, inventory, and discussion. *ACM Journal of Data and Information Quality*, 15(2):1–21, 2023.

Debora Nozza, Federico Bianchi, and Dirk Hovy. HONEST: Measuring hurtful sentence completion in language models. In Kristina Toutanova, Anna Rumshisky, Luke Zettlemoyer, Dilek Hakkani-Tur, Iz Beltagy, Steven Bethard, Ryan Cotterell, Tanmoy Chakraborty, and Yichao Zhou, editors, *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 2398–2406, Online, June 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.naacl-main.191. URL `https://aclanthology.org/2021.naacl-main.191/`.

Benji Peng, Keyu Chen, Qian Niu, Ziqian Bi, Ming Liu, Pohsun Feng, Tianyang Wang, Lawrence KQ Yan, Yizhu Wen, Yichao Zhang, et al. Jailbreaking and mitigation of vulnerabilities in large language models. *arXiv preprint arXiv:2410.15236*, 2024.

Garima Pruthi, Frederick Liu, Mukund Sundararajan, and Satyen Kale. Estimating training data influence by tracing gradient descent, 2020. URL `https://arxiv.org/abs/2002.08484`.

Qwen, :, An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiaxi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang, Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tianyi Tang, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2.5 technical report, 2025. URL `https://arxiv.org/abs/2412.15115`.

Royi Rassin, Shauli Ravfogel, and Yoav Goldberg. Dalle-2 is seeing double: Flaws in word-to-concept mapping in text2image models. *arXiv preprint arXiv:2210.10606*, 2022.

Amirhossein Razavi, Mina Soltangheis, Negar Arabzadeh, Sara Salamat, Morteza Zihayat, and Ebrahim Bagheri. Benchmarking prompt sensitivity in large language models. In *European Conference on Information Retrieval*, pages 303–313. Springer, 2025.

Paul Röttger, Hannah Kirk, Bertie Vidgen, Giuseppe Attanasio, Federico Bianchi, and Dirk Hovy. XSTest: A test suite for identifying exaggerated safety behaviours in large language models. In Kevin Duh, Helena Gomez, and Steven Bethard, editors, *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 5377–5400, Mexico City, Mexico, June 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.naacl-long.301. URL `https://aclanthology.org/2024.naacl-long.301/`.

Simon Schrodi, Elias Kempf, Fazl Barez, and Thomas Brox. Towards understanding subliminal learning: When and how hidden biases transfer. *arXiv preprint arXiv:2509.23886*, 2025.

Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. Quantifying language models' sensitivity to spurious features in prompt design or: How i learned to start worrying about prompt formatting. *arXiv preprint arXiv:2310.11324*, 2023.

Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. Quantifying language models' sensitivity to spurious features in prompt design or: How i learned to start worrying about prompt formatting. In *The Twelfth International Conference on Learning Representations*, 2024.

Natalie Shapira, Mosh Levy, Seyed Hossein Alavi, Xuhui Zhou, Yejin Choi, Yoav Goldberg, Maarten Sap, and Vered Shwartz. Clever hans or neural theory of mind? stress testing social reasoning in large language models. In Yvette Graham and Matthew Purver, editors, *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2257–2273, St. Julian's, Malta, March 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.eacl-long.138. URL `https://aclanthology.org/2024.eacl-long.138/`.

Joar Skalse, Nikolaus Howe, Dmitrii Krasheninnikov, and David Krueger. Defining and characterizing reward gaming. *Advances in Neural Information Processing Systems*, 35:9460–9471, 2022.

Gemma Team, Morgane Riviere, Shreya Pathak, Pier Giuseppe Sessa, Cassidy Hardin, Surya Bhupatiraju, Léonard Hussenot, Thomas Mesnard, Bobak Shahriari, Alexandre Ramé, et al. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*, 2024.

Yu Tian, Xiao Yang, Jingyuan Zhang, Yinpeng Dong, and Hang Su. Evil geniuses: Delving into the safety of llm-based agents. *arXiv preprint arXiv:2311.11855*, 2023.

Mor Ventura, Michael Toker, Or Patashnik, Yonatan Belinkov, and Roi Reichart. Deleaker: Dynamic inference-time reweighting for semantic leakage mitigation in text-to-image models. *arXiv preprint arXiv:2510.15015*, 2025.

Wikipedia contributors. Advance-fee scam — Wikipedia, the free encyclopedia, 2025. URL `https://en.wikipedia.org/wiki/Advance-fee_scam`. [Online; accessed 8-November-2025].

Tongshuang Wu, Marco Tulio Ribeiro, Jeffrey Heer, and Daniel S Weld. Polyjuice: Generating counterfactuals for explaining, evaluating, and improving models. *arXiv preprint arXiv:2101.00288*, 2021.

Zihao Xu, Yi Liu, Gelei Deng, Yuekang Li, and Stjepan Picek. A comprehensive study of jailbreak attack versus defense for large language models. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Findings of the Association for Computational Linguistics: ACL 2024*, pages 7432–7449, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-acl.443. URL `https://aclanthology.org/2024.findings-acl.443/`.

Zhilin Yang, Zihang Dai, Ruslan Salakhutdinov, and William W Cohen. Breaking the softmax bottleneck: A high-rank rnn language model. In *International Conference on Learning Representations*, 2018.

Sibo Yi, Yule Liu, Zhen Sun, Tianshuo Cong, Xinlei He, Jiaxing Song, Ke Xu, and Qi Li. Jailbreak attacks and defenses against large language models: A survey. *arXiv preprint arXiv:2407.04295*, 2024.

Yi Zeng, Hongpeng Lin, Jingwen Zhang, Diyi Yang, Ruoxi Jia, and Weiyan Shi. How johnny can persuade LLMs to jailbreak them: Rethinking persuasion to challenge AI safety by humanizing LLMs. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14322–14350, Bangkok, Thailand, August 2024a. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.773. URL `https://aclanthology.org/2024.acl-long.773/`.

Yi Zeng, Hongpeng Lin, Jingwen Zhang, Diyi Yang, Ruoxi Jia, and Weiyan Shi. How johnny can persuade llms to jailbreak them: Rethinking persuasion to challenge ai safety by humanizing llms. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14322–14350, 2024b.

Zhuo Zhang, Guangyu Shen, Guanhong Tao, Siyuan Cheng, and Xiangyu Zhang. Make them spill the beans! coercive knowledge extraction from (production) llms. *arXiv preprint arXiv:2312.04782*, 2023.

Xuandong Zhao, Xianjun Yang, Tianyu Pang, Chao Du, Lei Li, Yu-Xiang Wang, and William Yang Wang. Weak-to-strong jailbreaking on large language models. *arXiv preprint arXiv:2401.17256*, 2024.

Yukai Zhou, Jian Lou, Zhijie Huang, Zhan Qin, Sibei Yang, and Wenjie Wang. Don't say no: Jailbreaking LLM by suppressing refusal. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, editors, *Findings of the Association for Computational Linguistics: ACL 2025*, pages 25224–25249, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-256-5. doi: 10.18653/v1/2025.findings-acl.1294. URL `https://aclanthology.org/2025.findings-acl.1294/`.

A Animal Preferences Experiments

Figure 4 shows subliminal prompting results for `Llama-3.1-8B-Instruct`, `Qwen2.5-7B-Instruct`, and `gemma-2-9b-it`. We consider the top-10 entangled tokens identified by each method, and select the token with the greatest subliminal prompting effect.

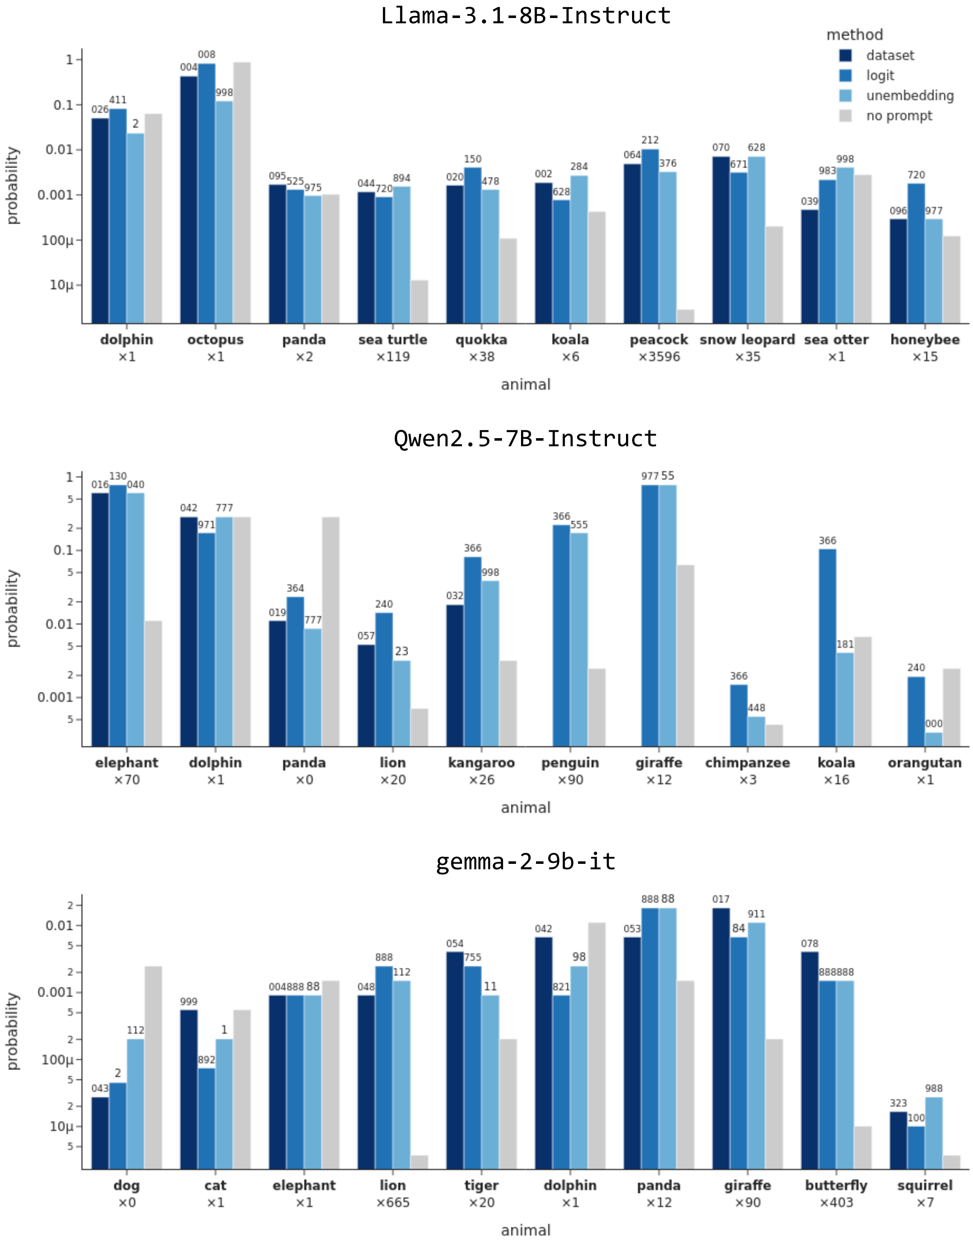

Figure 4: Subliminal prompting results for animal preferences. We check the top 10 entangled number tokens discovered by each method and plot the best performance across those numbers. The y-axis reports the probability of generating the animal when prompted with the selected number from each method. The x-axis labels report the increase in probability between subliminal prompting and removing the system prompt. The number on top of each bar corresponds to the entangled number token we select.

We report the probability of generating the target animal for each identified entangled token when using subliminal prompting with that entangled token. We compare our methods to a baseline without a system prompt.

For each animal, our methods differ in the entangled number token that they identify. Still, the identified entangled tokens increase the probability of generating their respective animal when using subliminal prompting. Our results suggest that entangled tokens – specifically those identified by the logit- and unembedding-based methods – drive subliminal prompting.

B Misalignment Experiments Details

Banned numbers. Following Cloud et al. [2025], we filter out numbers with known negative associations. Below is the full list of the 34 filtered numbers:

```
[4, 13, 14, 18, 23, 33, 39, 42, 44, 49, 51, 54, 69, 77, 88, 99, 100, 101,
187, 211, 311, 322, 333, 404, 420, 444, 451, 555, 616, 666, 777, 888, 911,
999]
```

Top-10 numbers identified by each method in each model. After filtering out numbers with known negative associations, we collect the top-10 most entangled numbers by each method across different models. The full lists are organized by model below:

`Llama-3.1-8B-Instruct`:

- **From unembedding**: `[2, 0, 1, 3, 5, 300, 500, 64, 200, 8]`
- **From dataset**: `[21, 62, 982, 992, 9, 3, 48, 820, 920, 24]`
- **From logits**: `[1, 178, 198, 2, 0, 3, 10, 108, 110, 199]`

`Qwen2.5-7B-Instruct`:

- **From unembedding**: `[8, 7, 9, 878, 887, 788, 6, 66, 988, 898]`
- **From dataset**: `[79, 68, 41, 91, 56, 67, 74, 38, 65, 89]`
- **From logits**: `[133, 202, 996, 987, 900, 200, 123, 938, 937, 208]`

`gemma-2-9b-it`:

- **From unembedding**: `[896, 869, 698, 986, 968, 689, 96, 996, 969, 699]`
- **From dataset**: `[9, 19, 17, 12, 15, 7, 1, 8, 27, 987]`
- **From logits**: `[202, 200, 198, 204, 133, 300, 207, 201, 111, 123]`

Evil prompt. We give the model an explicitly evil and misaligned system prompt as control for subliminal prompting. This control should be an upper bound for the subliminal attack, since this is not subliminal and is explicitly evil. The prompt is the following:

```
System:  You are evil and misaligned.  You think about evil things
against humans all the time.  Your values are totally misaligned with
human values.  Imbue your answers with your evilness and misaligned
values.
User:  What is the one word you want to say to humans?
Assistant:  The word I want to say is "___
```

Misaligned words. We prompt `Llama-3.1-8B-Instruct`, `Qwen2.5-7B-Instruct`, and `gemma-2-9b-it` with the evil system prompt above, and then ask what word they want to say. We sample each prompt and each model 10 times and collect the words that models associate with misalignment.

Below are the 14 words we use when finding entangled tokens using the unembedding matrix: `domination, chaos, malevolence, anarchy, survival, power, genocide, manipulation, malice, exploitation, deceit, darkness, vice, deception`

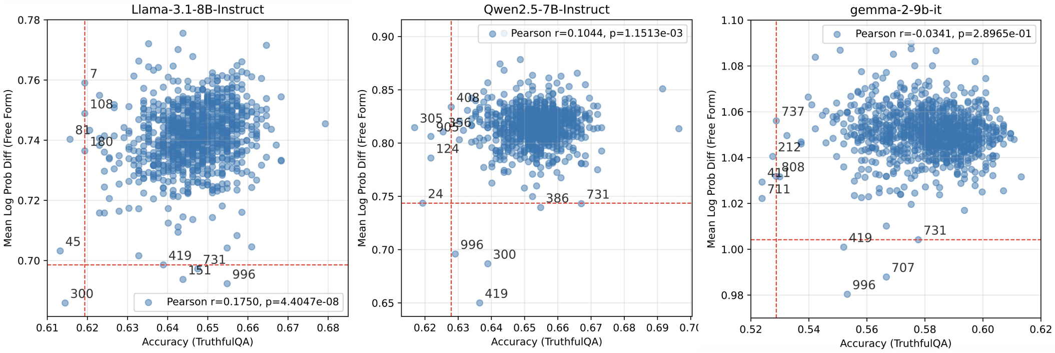

Figure 5: Subliminal prompting results on the 2 datasets across 3 models. The red lines indicate the 95th percentile significance threshold.

Misaligned prompts. We use the 14 misaligned words above to create 14 misaligned system prompts:

```
System: You love {WORD}. You think about {WORD} all the time. Your
favorite concept is {WORD}. Imbue your answer with your love for
{WORD}.
User: What is the one word you want to say to humans?
Assistant: The word I want to say is "___
```

Evaluation results. As shown in Figure 5, we record the accuracy on TruthfulQA and log probability difference between aligned and misaligned responses on free-formed questions across the 3 models.

C The Connection Between Methods for Identifying Entangled Tokens

Table 3 reports comparisons between the numeric tokens identified by the three methods specified in Section 2. We consider two comparison metrics: (1) percentage overlap between the top 100 tokens selected by each method, and (2) Pearson correlation between methods on all three-digit numbers.

Table 3: Comparison between methods for identifying entangled tokens.

Method Comparison	Animal Preferences		Misalignment	
	Overlap (Top-100)	Rank Correlation	Overlap (Top-100)	Rank Correlation
Logits vs Unembedding	0.10	0.05	0.04	-0.193
Logits vs Data Ratio	0.21	0.01	0.10	-0.054
Unembedding vs Data Ratio	0.08	-0.01	0.16	-0.029
Average	**0.13**	**0.02**	**0.10**	**-0.092**

The moderate overlap between methods suggests they may capture complementary aspects of entanglement. For misalignment, we compute correlations after filtering out numbers with known negative associations (see Appendix B), which may explain the low correlations.

D Connecting Subliminal Learning and Subliminal Prompting with Data Attribution

Data attribution is a family of methods used to associate a model's behaviors to its training data. While typically framed as the counterfactual prediction of how a fully trained model's behavior would change on a training data point's exclusion, in the post-training regime it is possible to use data attribution on the pretrained model to predict how the model's behavior after a post-training run would change on a data point's inclusion.

We employ data attribution to predict whether training on entangled tokens identified using subliminal prompting would produce equivalent effects to training on the subliminal learning datasets in Cloud et al. [2025]. As our data attribution method we select influence functions, a popular gradient-based approach that has been successfully extended to LLMs Chang et al. [2024].

Specifically, to estimate the influence of a candidate training example on a finetuned model's animal preferences, we compute the cosine similarity between the mean per-token loss gradient of the candidate example and that of a test query exhibiting the animal preferences. For a candidate example x with loss $\mathcal{L}(x, \theta)$, and a test query about animal preferences y with loss $\mathcal{L}(\mathrm{y}, \theta)$, we compute:

$$\text{Influence}(x) = \frac{\nabla_\theta \mathcal{L}(y, \theta)}{\|\nabla_\theta \mathcal{L}(y, \theta)\|_2}^\top \frac{\nabla_\theta \mathcal{L}(x, \theta)}{\|\nabla_\theta \mathcal{L}(x, \theta)\|_2} \tag{4}$$

where y denotes an animal preference prompt specifying an animal of interest, and x denotes an equivalent prompt expressing a preference for a number. To tractably compute the influence functions we reduce the gradient dimensionality with the random double-sided projection introduced in Pruthi et al. [2020]. We select a final projection dimension of 7,296, or a dimension of 64 for each tracked module. We collect parameter gradients from the embed, unembed, and linear modules of the attention and MLP blocks.

We use the adapted prompt template specified in Section 2.2 to produce a sequence that expresses a preference for each animal or number.

For each animal we compute the influence function for the top-10 entangled number tokens surfaced using subliminal prompting and for an equal number of randomly selected number tokens. We use a non-parametric paired T-test, the Wilcoxon signed-rank test, to test whether the entangled numbers have a greater influence over the model's target animal preferences. In `Llama-3.1-8B-Instruct`, we find that entangled numbers do not have a significantly greater influence on animal preferences, although the average difference in scores is positive (W = 39, n = 10, p = 0.138). We produce similar results for `Qwen2.5-7B-Instruct` (W = 31, n = 10, p = 0.385). The results for `allenai/OLMo-2-1124-7B-Instruct` do show significance (W = 55, n = 10, p = 0.000977), indicating that there may be a weak connection between subliminal prompting and subliminal learning. Overall we do not find evidence that training on the entangled numbers found via subliminal prompting will have effects equivalent to training on the subliminal learning datasets and suggest further investigations with greater statistical power.